

# **CS208: Applied Privacy for Data Science** **Machine Learning under DP**

School of Engineering & Applied Sciences  
Harvard University

March 8, 2022

# ML Inputs and Loss Functions


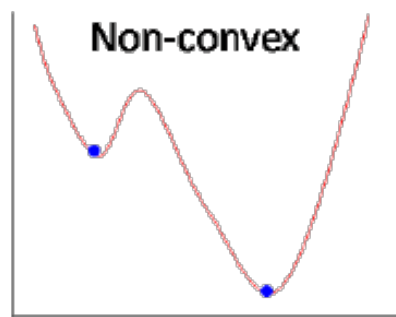
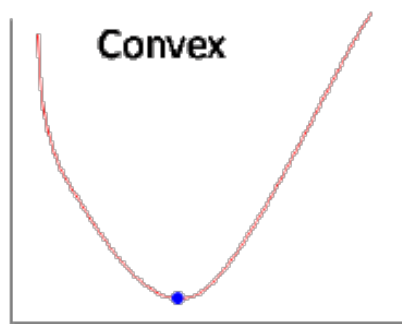
- **Data:**  $(x_1, y_1), \dots, (x_n, y_n) \sim \mathcal{P}$ 
  - Examples  $x_i \in \mathcal{X}$   $d$ -dimensional, discrete or continuous
  - Labels  $y_i \in \mathcal{Y}$  1-dimensional, discrete or continuous
  - $\mathcal{P}$  typically unknown
- **A loss function:**
  - $\ell : \Theta \times \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$      $\ell(\theta|x_i, y_i)$  measures ``loss"
  - Define  $L: \Theta \rightarrow \mathbb{R}$      $L(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(\theta|x_i, y_i)$
  - Example: Squared loss:  $\ell(\theta|x, y) = |(\beta_1 x + \beta_0) - y|^2$ .
- **Goal:** output  $\hat{\theta} \in \Theta$  s.t.  
$$L(\hat{\theta}) \approx \min L(\theta)$$

# Convexity

- **Def:**  $L$  is **convex** if for all points  $\vec{a}, \vec{b}$ , we have

$$L\left(\frac{\vec{a} + \vec{b}}{2}\right) \leq \frac{L(\vec{a}) + L(\vec{b})}{2}.$$



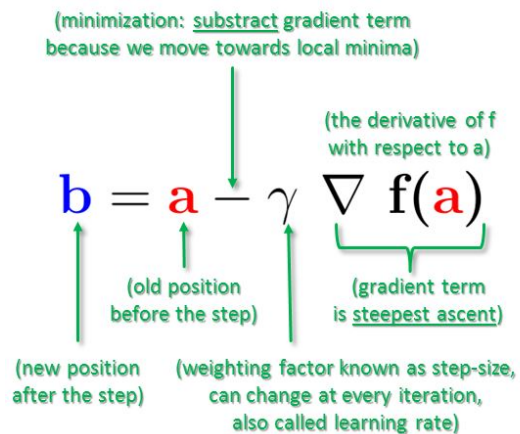
- Convex functions have **no local minima**



- Loss function for logistic regression is convex
  - No closed form solution for minimum, but it is easy to find

# Gradient Descent

- Proceed in steps
- Start from (carefully chosen) initial parameters  $\hat{\theta}_0$
- At each step, move in direction opposite to the gradient of the loss  $\nabla L(\hat{\theta} \mid \vec{x}, \vec{y})$ .
- Gradient is the vector of partial derivatives



# Gradient Descent

- Specify
  - Number of steps  $T$
  - Learning rate  $\eta$
- Pick initial point  $\hat{\theta}_0 \in \Theta$
- For  $t = 1$  to  $T$ 
  - Compute gradient

$$g_t = \nabla L(\hat{\theta}_{t-1}) = \frac{1}{n} \sum_i \nabla \ell(\hat{\theta}_{t-1} | x_i, y_i)$$

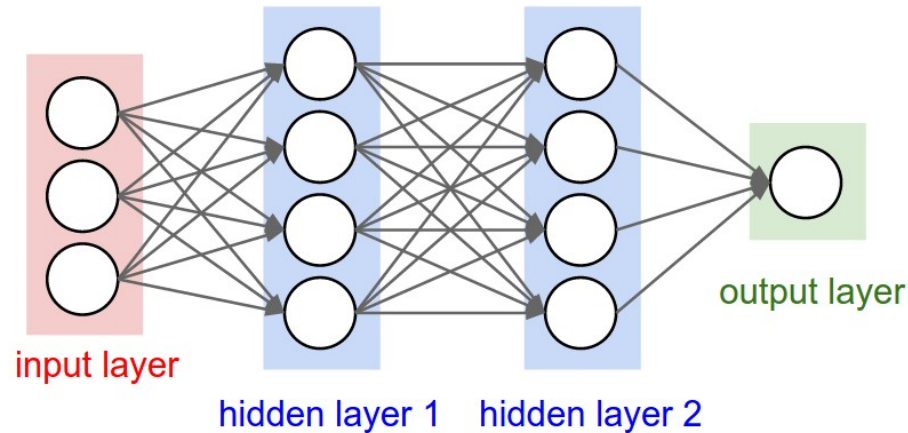
- $\hat{\theta}_t = \hat{\theta}_{t-1} - \eta \cdot g_t$

- Output  $\hat{\theta} = \sum_{t=1}^T \hat{\theta}_t$  or  $\hat{\theta}_T$

Average iterate


Last iterate

# Gradient Descent for Neural Networks

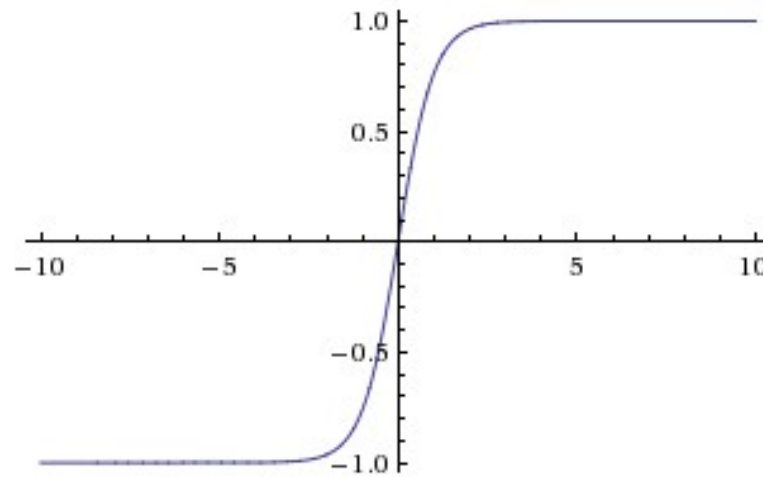


- Each node is a linear function of inputs (specified by  $\theta$ ) composed with a nonlinear “activation” function
- Gradient of Loss function can be computed quickly
  - Using chain rule (technique called “backpropagation”)
- But no longer convex, has many local minima
  - Can get stuck in a bad place
  - But works well in practice!

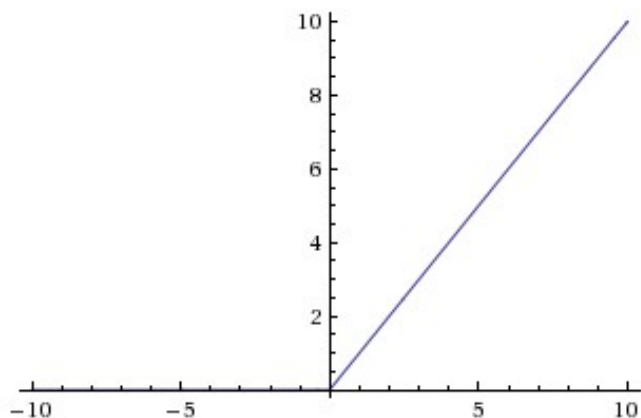
# Common Activation Functions



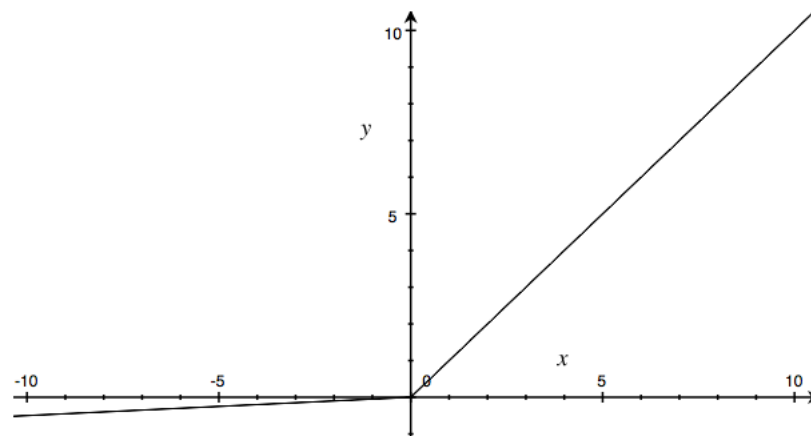
$$\text{Sigmoid } \sigma(x) = \frac{1}{1+e^{-x}}$$



$$\tanh(x) = 2\sigma(2x) - 1$$



$$\text{ReLU}(x) = \max(0, x)$$



$$\text{Leaky ReLU}(x) = \max(0.05x, x)$$

# DP for Vector-Valued Functions

- Let  $f : \mathcal{X}^n \rightarrow \mathbb{R}^k$ , and  $M(x) = f(x) + Z$  for noise  $Z \in \mathbb{R}^k$ .
- global  $\ell_2$ -sensitivity of  $f$  is

$$\text{GS}_{f, \ell_2} \stackrel{\text{def}}{=} \max_{x \sim x'} \|f(x) - f(x')\|_2.$$

$$\|z\|_2 = \left( \sum_j |z_j|^2 \right)^{1/2}$$

- **Gaussian Mechanism:**  $Z \sim \mathcal{N}\left(\vec{0}, 2 \left(\frac{\text{GS}_{f, \ell_2}}{\varepsilon}\right)^2 \cdot \ln \frac{1.25}{\delta} \cdot I_k\right)$ 
  - independent Gaussian noise per coordinate.



# Robustness to Noise in Gradient Estimation

- For efficiency:

Sample a minibatch  $B \in \{1, 2, \dots, n\}$

Gradient estimate  $\tilde{g}_t = \frac{1}{|B|} \sum_{i \in B} \nabla \ell(\hat{\theta}_{t-1}, x_i, y_i)$

Stochastic Gradient Descent (SGD)!

- For privacy:

Add Gaussian Noise  $\tilde{g}_t = g_t + \mathcal{N}(0, \sigma^2 I)$

In both cases,  $\tilde{g}_t$  is an unbiased estimate of  $g_t$ :  $E[\tilde{g}_t] = g_t$

# DP Gradient Descent

[Williams-McSherry'10, ...]

- Specify
  - Number of steps  $T$
  - Learning rate  $\eta$
  - Privacy parameters  $\epsilon, \delta$
  - Clipping parameter  $C$ . Write  $[\vec{z}]_{\Delta} = \vec{z} \cdot \max\left(1, \frac{C}{\|\vec{z}\|_2}\right)$ .
  - Noise variance  $\sigma^2 = \text{TBD}(T, \epsilon, \delta, C)$ .
- Pick initial point  $\hat{\theta}_0$
- For  $t = 1$  to  $T$ 
  - Estimate gradient as **noisy** average of **clipped** gradients
$$\hat{g}_t = \frac{1}{n} \sum_i [\nabla \ell(\hat{\theta}_{t-1} | x_i, y_i)]_C + \mathcal{N}(0, \sigma^2 I)$$
  - $\hat{\theta}_t = \hat{\theta}_{t-1} - \eta \cdot \hat{g}_t$
- Output  $\hat{\theta} = \sum_{t=1}^T \hat{\theta}_t$  or  $\hat{\theta}_T$

# Privacy Analysis

- Proof idea: Show releasing  $(\hat{g}_1, \hat{g}_2, \dots, \hat{g}_T)$  satisfies DP
  - Each step (releasing  $\hat{g}_t$ ) satisfies  $(\epsilon, \delta)$ -DP
  - Adaptive composition across  $T$  steps

# Privacy Analysis

- By Gaussian Mechanism, each iteration is  $(\epsilon_0, \delta_0)$ -DP if

$$\sigma^2 = 2 \left( \frac{C}{\epsilon_0 n} \right)^2 \cdot \ln \frac{1.25}{\delta_0}$$

- By Advanced Composition for **adaptive** queries, overall algorithm is  $(\epsilon, \delta)$ -DP for:

$$\begin{aligned} \epsilon &= O \left( \epsilon_0 \cdot \sqrt{T \ln(2/\delta)} \right) \\ \delta &= 2T \cdot \delta_0 \end{aligned}$$

- Solving, suffices to use noise variance

$$\sigma^2 = O \left( \left( \frac{C}{\epsilon n} \right)^2 \cdot T \cdot \ln \frac{T}{\delta} \cdot \ln \frac{1}{\delta} \right)$$

# Improved Analysis with “Concentrated DP”

[Dwork-Rothblum '16, Bun-Steinke '16]

- By Gaussian Mechanism, each iteration is  $\varepsilon_0^2$ -zCDP if

$$\sigma^2 = \frac{1}{2} \left( \frac{C}{\varepsilon_0 n} \right)^2 \cdot \ln \frac{1.25}{\delta_0}$$

- By composition of zCDP, overall algorithm is  $T \cdot \varepsilon_0^2$ -zCDP.
- By properties of zCDP, overall algorithm is  $(\varepsilon, \delta)$ -DP for:

$$\varepsilon = T \cdot \varepsilon_0^2 + 2 \sqrt{T \cdot \varepsilon_0^2 \cdot \ln(1/\delta)}$$

- Solving, suffices to use noise variance

$$\sigma^2 = O \left( \left( \frac{C}{\varepsilon n} \right)^2 \cdot T \cdot \ln \frac{1}{\delta} \cdot \ln \frac{T}{\delta} \right)$$

# DP **Stochastic** Gradient Descent (**S**GD)

[Jain-Kothari-Thakurta '12, Song-Chaudhuri-Sarwate '13, Bassily-Smith-Thakurta '14]

- Specify
  - Number of steps  $T$ , learning rate  $\eta$ , privacy parameters  $\varepsilon, \delta$ , clipping parameter  $C$ .
  - Batch size  $B \ll n$  (for efficiency)
  - Noise variance  $\sigma^2 = \text{TBD}(T, \varepsilon, \delta, C, B)$ .
- Pick initial point  $\hat{\theta}_0$
- For  $t = 1$  to  $T$ 
  - Select a random batch  $S_t \subseteq \{1, \dots, n\}$  of size  $B$ .
  - Estimate gradient as noisy average of clipped gradients
$$\hat{g}_t = \frac{1}{B} \sum_{i \in S_t} [\nabla \ell(\hat{\theta}_{t-1} | x_i, y_i)]_C + \mathcal{N}(0, \sigma^2 I)$$
  - $\hat{\theta}_t = \hat{\theta}_{t-1} - \eta \cdot \hat{g}_t$
- Output  $\hat{\theta} = \sum_{t=1}^T \hat{\theta}_t$  or  $\hat{\theta}_T$

# DP SGD: Improved Privacy Analysis

[Bassily-Smith-Thakurta '14, Abadi-Chu-Goodfellow-McMahan-Mironov-Talwar-Zhang '17]

- **Privacy amplification by subsampling:**

If  $S : \mathcal{X}^n \rightarrow \mathcal{X}^B$  outputs a random subset of  $p$  out of  $n$  rows and  $M : \mathcal{X}^B \rightarrow \mathcal{Y}$  is  $(\epsilon, \delta)$ -DP, then

$M'(x) = M(S(x))$  is  $(\ln(1 + (e^\epsilon - 1)p), p\delta)$ -DP.

– Keep  $S_t$  secret; Use their randomness



$p\epsilon$

- Poisson sampling: choosing each point independently with probability  $p=B/n$ .
- Choosing  $B$  points without replacement
- Choosing  $B$  points with replacement

# DP SGD: Improved Privacy Analysis

[Bassily-Smith-Thakurta '14, Abadi-Chu-Goodfellow-McMahan-Mironov-Talwar-Zhang '17]

- **Privacy amplification by subsampling:**

If  $S : \mathcal{X}^n \rightarrow \mathcal{X}^B$  outputs a random subset of  $pn$  out of  $n$  rows and  $M : \mathcal{X}^B \rightarrow \mathcal{Y}$  is  $(\epsilon, \delta)$ -DP, then

$M'(x) = M(S(x))$  is  $(\ln(1 + (e^\epsilon - 1)p), p\delta)$ -DP.

- Keep  $S_t$  secret; Use their randomness



$p\epsilon$

- We can take  $p = B/n$ .

- Unfortunately privacy amplification by subsampling does not hold for zCDP.
- But similar analysis can be recovered using the “moments accountant” [Abadi et al. '17] or “truncated zCDP” [Bun et al. '18].



# DP SGD: Privacy Analysis

- By Gaussian Mechanism, each iteration is  $(\epsilon_0, \delta_0)$ -DP if

$$\sigma^2 = 2 \left( \frac{C}{\epsilon_0 n} \right)^2 \cdot \ln \frac{1.25}{\delta_0}$$

- Subsampling + Gaussian is  $(\epsilon'_0, \delta'_0)$ -DP

$$\epsilon'_0 = \frac{B}{n} \epsilon_0 \quad \delta'_0 = \frac{B}{n} \delta_0$$

- Advanced Composition over  $T$  iterations

$$\begin{aligned} \epsilon &= O \left( \epsilon'_0 \cdot \sqrt{T \ln(2/\delta)} \right) \\ \delta &= 2T \cdot \delta'_0 \end{aligned}$$

Moments  
accountant  
 $O(\frac{B}{n} \epsilon_0 \sqrt{T}, T\delta)$ ;  
[WBK19,MTZ19]  
better analysis;  
[JUO20] attack

Main Idea: Privacy amplification by subsampling + Composition

# Neural Networks & Privacy

- Choice of the model architecture
  - The Gaussian noise proportional to the square root of number of parameters.
- Hyperparameter tuning
  - If we run analyses on the training data with various hyperparameter settings, and choose the best one (any problems?)
  - Doing this privately (with additional cost in privacy)
  - Use public dataset (CIFAR-100 dataset for CIFAR-10 training)
- Still can be improved...
  - MNIST(99.8% baseline) 98.1%  $(2.93, 10^{-5})$  – *DP*
  - CIFAR-10(99.7% baseline) 66.2%  $(7.53, 10^{-5})$  – *DP*

Results from [PTS+ 20]

# **Differentially Private Empirical Risk Minimization**

# Supervised ML Output

## Primary Goal (risk minimization):

- Find  $\theta \in \Theta$  minimizing  $L(\theta) = \mathbb{E}_{(x,y) \sim \mathcal{P}}[\ell(\theta|x, y)]$ .
- Difficulty:  $\mathcal{P}$  unknown.

## Subgoal 1 (empirical risk minimization (ERM)):

- Find  $\theta \in \Theta$  minimizing  $L(\theta|\vec{x}, \vec{y}) = \frac{1}{n} \sum_{i=1}^n \ell(\theta|x_i, y_i)$ .
- Turns learning into optimization.
- Difficulty: overfitting\*

## Subgoal 2 (regularized ERM):

- Find  $\theta \in \Theta$  minimizing  $L(\theta|\vec{x}, \vec{y}) = \frac{1}{n} \sum_{i=1}^n \ell(\theta|x_i, y_i) + R(\theta)$ .
- $R(\theta)$  typically penalizes “large”  $\theta$ , can capture Bayesian prior.

\*Fact: DP automatically helps prevent overfitting! [Dwork et al. '15]

# Output Perturbation

[Chaudhuri-Monteleoni-Sarwate '11]

$$M(\vec{x}, \vec{y}) = \operatorname{argmin}_{\theta} \left( \frac{1}{n} \sum_{i=1}^n \ell(\theta | x_i, y_i) + R(\theta) \right) + \text{Noise}$$

**Challenge:** bounding sensitivity of  $\theta_{opt} = \operatorname{argmin}_{\theta}(\cdot)$

- Global sensitivity can be infinite (e.g. OLS regression)
- Global sensitivity can be bounded when  $\ell$  is strictly convex, has bounded gradient (as a function of  $\theta$ ), and  $R$  is strongly convex. Even analyzing local sensitivity seems to require unique global optimum and using an optimizer that is guaranteed to succeed.

# Objective Perturbation

[Chaudhuri-Monteleoni-Sarwate '11]

$$M(\vec{x}, \vec{y}) = \operatorname{argmin}_{\theta} \left( \frac{1}{n} \sum_{i=1}^n \ell(\theta | x_i, y_i) + R(\theta) + R_{\text{priv}}(\theta, \text{noise}) \right)$$

**Challenge:** how to put noise in the objective function?

- [CMS11] use  $R_{\text{priv}}(\theta, v) = \langle \theta, v \rangle + c \|\theta\|^2$  where  $v$  is sampled with probability density  $\propto \exp(-c' \varepsilon \|v\|)$ .
- Privacy proven under similar assumptions on  $\ell$  and  $R$  as before, plus  $\ell$  having bounded Jacobian.
- Has better performance than output perturbation [CMS11].

# Exponential Mechanism for ML

[Kasiwiswanathan-Lee-Nissim-Raskhodnikova-Smith '11]

Use utility function

$$u((\vec{x}, \vec{y}), \theta) = -L(\theta|\vec{x}, \vec{y}) = -\frac{1}{n} \sum_{i=1}^n \ell(\theta|x_i, y_i) - R(\theta).$$

That is,

$$\Pr[M(\vec{x}, \vec{y}) = \theta] \propto e^{-\frac{\varepsilon}{2} \sum_{i=1}^n \ell(\theta|x_i, y_i) - \frac{\varepsilon n}{2} R(\theta)}.$$

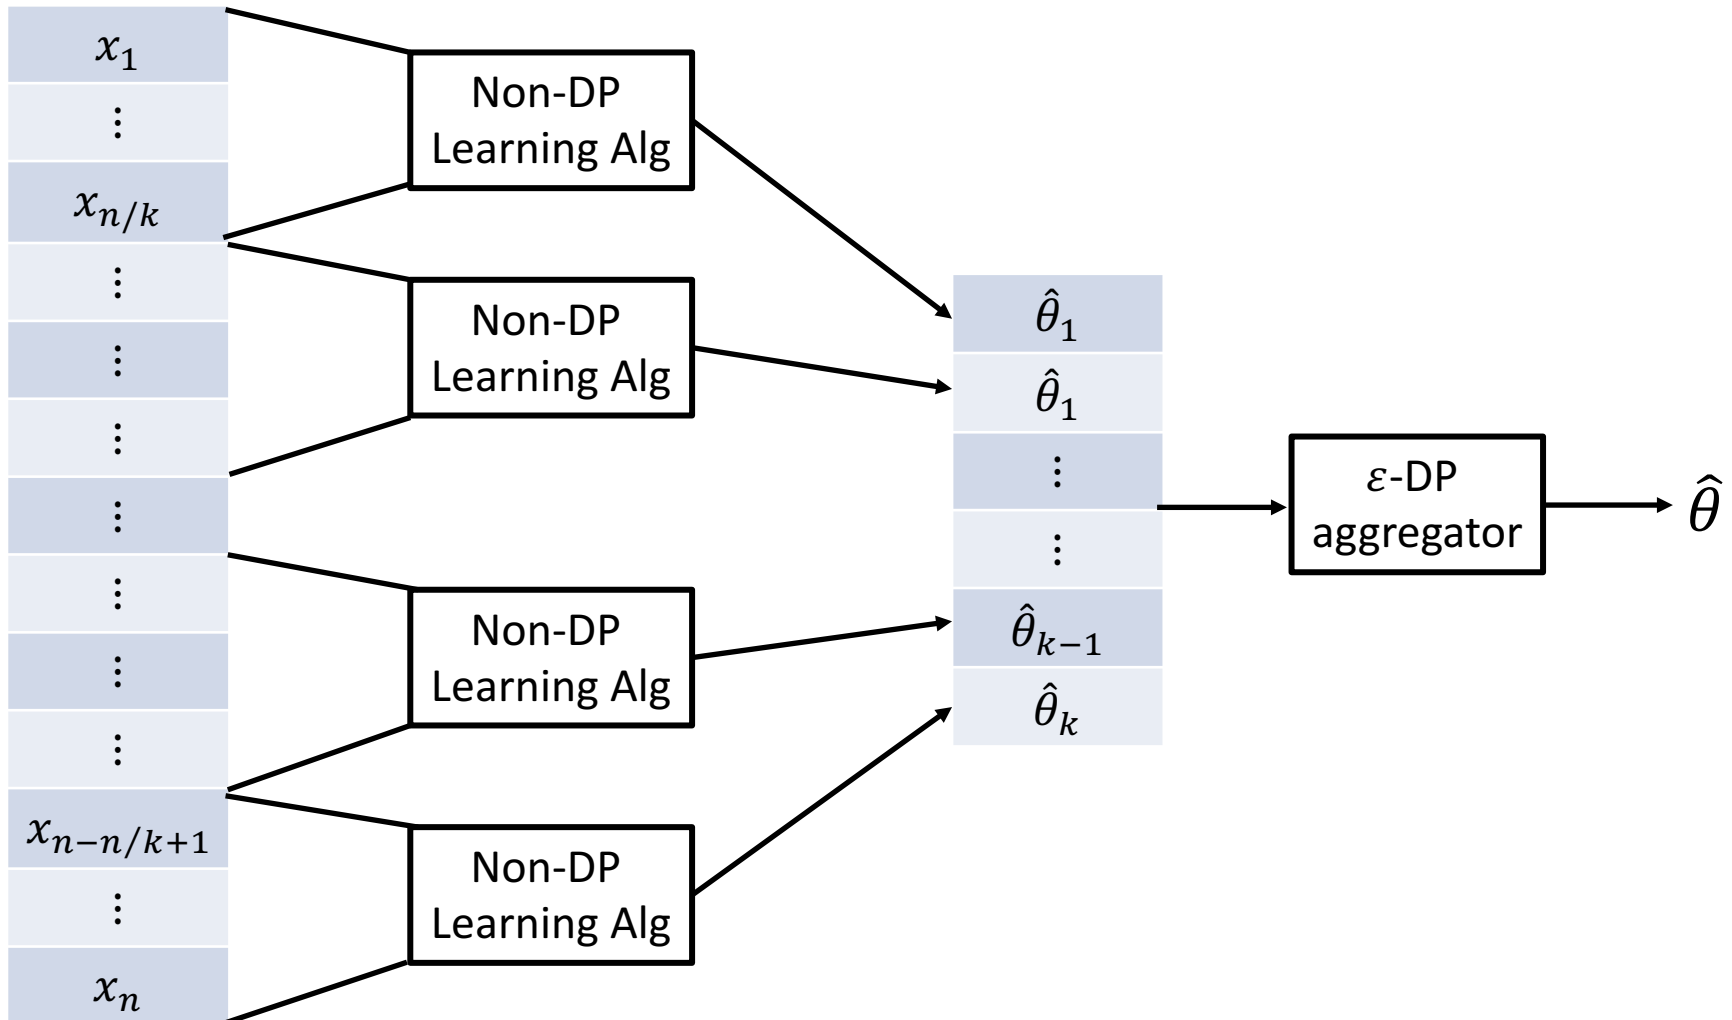
Is  $\varepsilon$ -DP if the loss functions are clipped to  $[0,1]$ . (why?)

**Thm [KLNRS '11, informally stated]:** anything learnable non-privately on a finite data universe is also learnable with DP (with larger  $n$ ).

**Problem:** runtime often exponential in dimensionality of  $\theta$ .

# Subsample & Aggregate

[Nissim-Rakhodnikova-Smith '07, Smith '11]



Q: Why is this  $\epsilon$ -DP?



# Subsample & Aggregate

[Nissim-Rakhodnikova-Smith '07, Smith '11]

- Typical aggregators: DP (clipped) mean, DP median
- **Benefits:**
  - Use any non-private estimator as a black box
  - Can give optimal asymptotic convergence rates: for many statistical estimators, variance is asymptotically  $c_\theta / (\text{sample size})$ , so variance of DP mean  $\hat{\theta}$  is
$$(1/k) \cdot (c_\theta \cdot k/n) + O(1/\varepsilon k)^2 = (1 + o(1)) \cdot c_\theta / n$$
if  $k = \omega(\sqrt{n})$ .
- **Drawbacks:**
  - Dependence on dimension, model parameters, distribution can be bad.
  - Often takes very large sample size to kick in.
- **Develop:**
  - Private Aggregation of Teacher Ensembles (PATE) [PAE+17, PSM+18]

# Modifying ML Algorithms

- **Another approach:** decompose existing ML/inference algorithms into steps that can be made DP, like Statistical Queries (estimating means of bounded functions)
- **Example:** linear regression
  - $S_{xx}/n, S_{xy}/n, \bar{x}, \bar{y}$  are all statistical queries